

Transactional Memory Evaluation using Apache Webserver

Haggai Eran

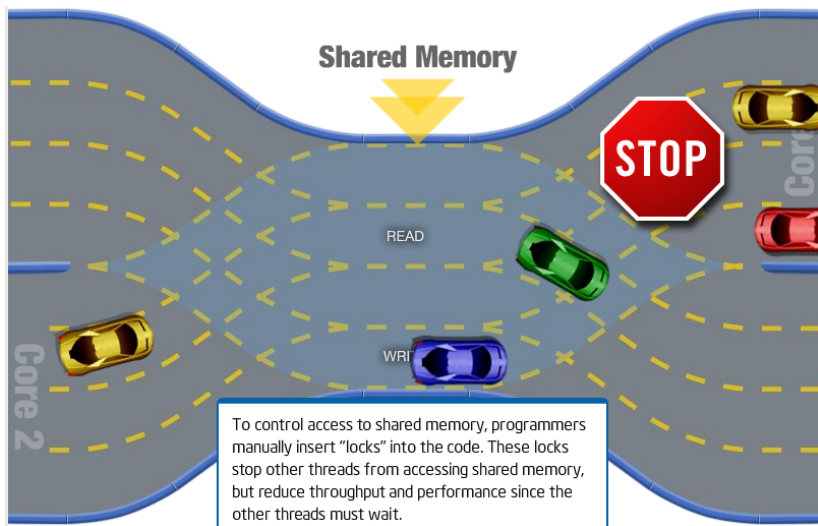
December 2, 2008

Traditional Synchronization


Example

```
void withdraw(account, amount) {  
    accounts[account] -= amount;  
}
```

Course-Grained Locks



Course-Grained Locks

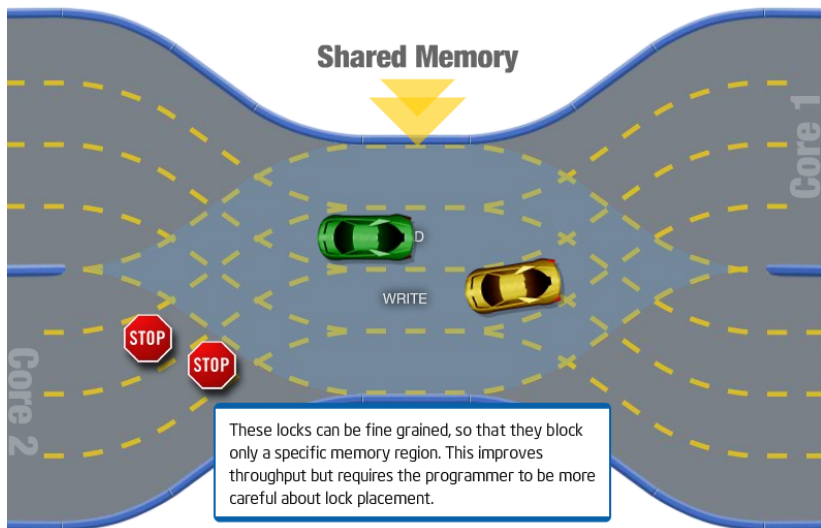


Example


```
void withdraw(account, amount) {  
    lock(big_mutex);  
    accounts[account] -= amount;  
    release(big_mutex);  
}
```

- Easy to program.
- Doesn't scale.

Fine-Grained Locks



Fine-Grained Locks




Example

```
void withdraw(account, amount) {  
    lock(accounts[account].mutex);  
    accounts[account] -= amount;  
    release(accounts[account].mutex);  
}
```

- Can scale well.
- Difficult to program.

Fine-Grained Locks Difficulties

Composition



Example

```
void transfer(fromAccount, toAccount, amount) {  
    withdraw(fromAccount, amount);  
    deposit(toAccount, amount);  
}
```

- Locking both accounts from transfer - breaks encapsulation, deadlocks.

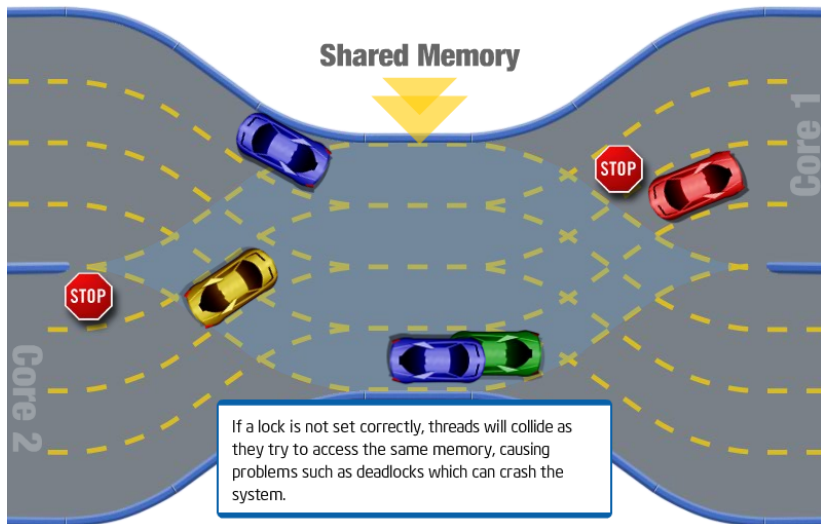
Fine-Grained Locks Difficulties

Locking Policies

Comment from the linux kernel

```
/*  
 * When a locked buffer is visible to the I/O layer  
 * BH_Laundry is set. This means before unlocking  
 * we must clear BH_Laundry, mb() on alpha and then  
 * clear BH_Lock, so no reader can see BH_Laundry set  
 * on an unlocked buffer and then risk to deadlock.  
 */
```


Fine-Grained Locks Difficulties



Transactional Memory

- Provide a simple API for programmers.
- Offering fast implementations.

Transactional Memory

Simple API

Example

```
void withdraw(account, amount) {  
    atomic {  
        accounts[account] -= amount;  
    }  
}
```

Nested transactions

```
void transfer(fromAccount, toAccount, amount) {  
    atomic {  
        withdraw(fromAccount, amount);  
        deposit(toAccount, amount);  
    }  
}
```


Transactional Memory

Implementation

- A transaction is run speculatively without taking any locks.
- Collisions are detected either at commit time or during the run.
- On collision, one of the transactions is aborted and its changes are rolled back.
- Later the aborted transaction is restarted.

Transactional Memory

Implementation



Transactional Memory

Implementation by software

- All global memory accesses are handled by a special library.
- The library detects collisions and handles commits and aborts.

Transactional Memory

Implementation by hardware

- Reuse the cache coherency mechanism in multicore/multiprocessor machines.
- Requires special hardware.
- Limitations: Size and duration of transactions, context switches.

Existing Benchmarks

- Red-Black trees benchmarks
- STAMP benchmark suite.
 - Bayesian network learning
 - Gene sequencing
 - Network intrusion detection
 - K-means clustering
 - Maze routing
 - Graph kernels
 - Client/server travel reservation system
 - Delaunay mesh refinement

Our Project's Goal

Create a benchmark based on a real-world application for transactional memory.


Apache Web Server



- Written in C.
- Support many Multiprocessing Modules (MPMs): Parallel execution strategies.
- A mainly developed threaded MPM is the Worker MPM: Runs several processes, each running a fixed number of threads.

Apache Web Server

Worker MPM



Apache Cache Module - mod_mem_cache

- There isn't much interaction between the worker threads.
- The cache module enables worker threads of the same process to share cached pages in memory.
- Currently implemented with one big lock.

Software Transactional Memory in C/C++

Library based

- Transactions denoted by beginning and ending function calls.
- Require accessing global variables through library functions.

Compiler based

- Transactions are denoted by special syntax.
- Automatic conversion of memory access (transactification).

Compiler based STM

	Tanger	ICC Experimental STM Compiler
License	Open source	Licensed for personal use
Backend Compiler	LLVM	ICC
STM Library	Any / tinySTM	Intel ITM Library
Function Calls	Transactify All	tm_callable
Indirect Functions	Resolved, at run-time	
Library Functions	Limited, Tarifa	malloc, free, some stdlib.h

Commit handlers

A common pattern we found, missing in both Tanger and ICC.

Example

```
__tm_atomic {  
    if (--obj->refcount == 0) {  
        cache_remove(cache, obj);  
        cleanup_cache_object(obj);  
    }  
}
```

Commit handlers

Should be converted to:

Example

```
__tm_atomic {  
    if (--obj->refcount == 0) {  
        cache_remove(cache, obj);  
    }  
}  
if (obj->refcount == 0)  
    cleanup_cache_object(obj);
```

Commit handlers

It would be nice to have:

Example

```
__tm_atomic {  
    if (--obj->refcount == 0) {  
        cache_remove(cache, obj);  
        on_commit(&cleanup_cache_object, obj);  
    }  
}
```



Evaluation

Evaluation of a web server requires:

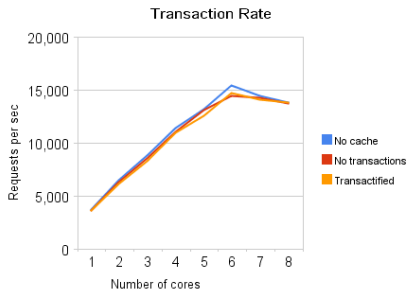
- A data set.
- Client strategy

We chose

- Data set of small files (man pages) so that the throughput of the NIC won't be the bottleneck.
- Running as many clients concurrently as possible to create contention on the server and its cache.
- Requesting pages according to Zipf distribution - to control locality.



Initial Results



Theory



- The linux file cache contains the entire data set => Apache's cache just gets in the way.
- Dynamically generated content might give the cache an advantage.

Computational Workload

- Serve the same workload though **man2html**, uncompressing and converting the files to HTML.



Computational Workload

$s = 1$




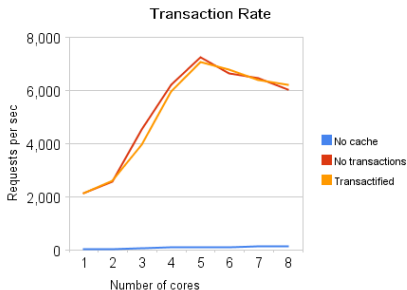
Computational Workload

$s = 2$



Computational Workload

$s = 3$



Thank you

Credits

- The car analogy graphics from Intel's presentation.
- The kernel quote from Maurice Herlihy's presentation.

Questions

?