

Transactifying Apache

Haggai Eran Ohad Lutzky Zvika Guz Idit Keidar

Department of Electrical Engineering, Technion - Israel Institute of Technology

haggaie@tx.technion.ac.il lutzky@gmail.com zguz@tx.technion.ac.il idish@ee.technion.ac.il

Abstract

Apache is a large-scale industrial multi-process and multi-threaded application, which uses lock-based synchronization. We report on our experience in modifying Apache to employ transactional memory instead of locks, a process we refer to as *transactification*; we are not aware of any previous efforts to transactify legacy software of such a large scale. Along the way, we learned some valuable lessons about which tools one should use, which parts of the code one should transactify and which are better left untouched, as well as on the intricacy of commit blocks. We also stumbled across weaknesses of existing software transactional memory (STM) toolkits, leading us to identify desirable features they are currently lacking. Finally, we present performance results from running Apache on a 32-core machine, showing that, surprisingly, there are scenarios where the performance of the STM-based version is close to that of the lock-based version. These results suggest that there are applications for which the overhead of using a software-only implementation of transactional memory is insignificant.

Categories and Subject Descriptors D.1.3 [Programming Techniques]: Concurrent Programming

General Terms Measurement, Performance, Experimentation

Keywords Software Transactional Memory

1. Introduction

The vast shift to multicore machines in the last couple of years indicates the emerging of the multicore era. In this context, Transactional Memory and Software Transactional Memory (STM) are of the primary tools targeted at allowing programmers to easily harness the parallelism of future multi-core machines and to extract the performance promise these systems endure.

One of the principal challenges all Transactional Memory system confront (beside delivering performance) is the ability to handle large-scale commercial applications and the ability to easily export large legacy code. As TM systems are maturing, the aspect of convertibility and completeness are becoming acute in order to allow TM to shift from concept into full-fledged commercial tool. In this work we try via a design example to answer how far from achieving the goal we are at.

To date, transactional memory was mostly employed within the niche of complex concurrent data structures, such as red-black trees and skip lists [9, 7], and isolated scientific algorithms (such as STAMP [3]). Moreover, it was mostly used to implement benchmarks explicitly with transactional memory from scratch. Additionally, we've seen benchmarks such as STMBench7 [8], which measures operations on a complex yet still artificial object graph.

In this paper, we use transactional memory for the first time in the context of large-scale industrial software, which, moreover, does not pertain to any of the typical niches of transactional memory benchmarks. Furthermore, we convert legacy code, which used lock-based synchronization, to work with transactional memory, rather than write the benchmark from scratch. We refer to this conversion process as *transactification*.

Our goal for this exercise is twofold. First, we wanted to examine whether the transactional memory

paradigm is broadly applicable, outside its traditional niches, in contexts where a high level of parallelism is already obtained using lock-based synchronization. [XX Can we quote something somebody said before on where it should not work?? XX].

Second, as transactional memory is being touted as the panacea for the limited parallelism of coarse-grained locks, and since coarse-grained locks that limit parallelism are widely used in legacy code, one can anticipate a growing need for transactification of legacy software. We sought to learn how painful such a process can be, and what can be done in order to facilitate it.

The rest of the paper is organized as follows: After a brief review of the apache application and the STM tools used in our work (Section 3), we describe in details the transactification process, including the hurdles we have encountered and the means used to overcome them (Section 4) including our wish-list (Section 5) for future STM versions. We then study (in Section 6) the performance of the transactified version, showing that the Transactified version is competitive to the legacy code in performance. This demonstrates that for some applications, the performance penalty of STM may not be as severe as usually assumed, and will hence justify the porting to TM with all the benefits it endures. Finally, we summarize (in Section 7) the lesson learned in the transactification process.

2. Related Work

3. Background: Software and Tools

3.1 Apache

Apache [16] HTTP server is a popular web server application written in C. It supports working on multiprocessor machines with several multi-processing modules (MPMs) each offering a different strategy for handling requests and distributing the work. The most popular threaded MPM is the *worker* MPM, which works by running multiple worker-threads under several processes, each thread handles a single request at a time. In each such process there are several worker threads, and also a listener thread that fetches incoming requests and dispatches them to the available workers, as illustrated in Figure 1.

There are not many points of interaction between the worker threads themselves, where transactional memory can be used. One such place is Apache's memory




Figure 1. Apache worker MPM architecture.

cache implemented by the `mod_mem_cache` module. This module enables the workers of each process to share a cache of recently served requests. A new request can be served from the memory cache, and save the time required to access the disk and generate the requested page. Since the cache is shared between multiple threads, it is synchronized by a single lock, therefore a good candidate for converting into transactional memory.

Apache's cache is implemented with a couple of modules. The first, `mod_cache`, implements the logic related to caching. It tests the metadata of each requests to see if it can be supplied from the cache, according to the request's HTTP headers and the system configuration. It uses one of the underlying cache implementation modules, `mod_mem_cache` or `mod_disk_cache` to do the actual caching.

The `mod_mem_cache` module implements a memory cache using a shared hash table and priority queue. The key to the hash table is the URL of the request, converted into a canonical form. The cache is limited both by size and by the number of elements, and by memory size, so on insertion, sometimes lower priority entries are removed from the cache. The priority is determined by one of two algorithms: LRU, removing the least recently used entries first, and GDSF (Greedy Dual Size Frequency) assigning score to entries based on the cost of a cache miss, and the entry size.

3.2 C STM Systems

C and C++ STM systems divide into two kinds: Library based and compiler based. Library based STMs [4, 6,

13] are built as a C library. Every transaction begins with a call into the library, and commits by another call. All reads and writes to global variables must be done through special library functions when in a transaction. This requires a great amount of work for converting an application to use STM. Not only accesses to global memory in the function that started the transaction must be converted, but also any access from any function being called from this function.

In contrast, compiler-based STM [5, 15, 1, 2] use a specialized compiler, which has extended syntax for transactional memory atomic blocks. The compiler can then automatically convert memory accesses inside transactions into calls to the underlying library. The term *transactification* sometimes refer to this process.

3.2.1 TANGER

The TANGER [5] transctifying compiler is an open-source academic compiler extension for LLVM [12], an extensible compiler framework. Tanger aims at creating a transactifying compiler that is independant of the STM system used. It works with the tinySTM [6] library, but can easily be extended to use other STM libraries by writing a simple plugin.

TANGER is accompanied with the TARIFA [5] tool, which transactifies compiled binaries, even without the sources, which might be a very important advantage when modifying legacy code.

3.2.2 Intel STM Compiler

Intel has published [15] an experimental STM compiler based on their industrial compiler ICC. It solved the above problem by adding some new function attributes to the language that tell the compiler which functions need to be transactified. The attribute `tm_callable` tells the compiler that a transactified version of the function will be needed. This way only functions that are required inside transactions can be marked as `tm_callable` and be transactified.

Although ICC uses a proprietary STM manager, Intel has published their ABI [10] allowing for other STM managers to replace their own. This feature and its selective transactification ability were the main reasons why we preferred ICC.

In the latest version of Intel STM Compiler, support for abort and commit handlers was in fact added to the system, by registering a callback function from inside a transaction.

An extension to the Gnu Compiler Collection (GCC) is being developed [1] to enable transactional memory support for GCC. It is intended to work with tinySTM, but being open source, other STM systems will probably be ported too. The syntax of the C/C++ language extensions is designed to be compatible with ICC. This means that applications converted for ICC will probably be compilable under GCC with this extension, without much modification.

4. Transactification

4.1 Which STM to use?

TANGER created a transactified version of each function in a compilation unit. Every function call inside a transaction was then converted to a call to the new version. This method is a major disadvantage when working on a large application. Many functions do not need a transactified version and this causes unneeded work for the compiler and the linker. Moreover, sometimes the transactification might fail because of calls to functions whose source is not available and cannot be transactified. This can cause the entire build process to fail, where in fact the code can be transactified without any error.

This was the main reason why we eventually chose to use Intel's STM Compiler. Lately a new release of tanger was announced, one in which the developer can annotate which functions should be transactified, however we didn't get a chance to try it.

4.2 What to transactify

We have decided to focus our change in Apache with the memory cache module. As described in Section 3.1, it is a significant point of interaction between Apache's worker threads, and since it is implemented with a course-grained lock, it is ideal for converting into STM. Another advantage is that it is a well encapsulated module, allowing us to see how can we convert only a small part of a larger system.

The conversion process included converting critical sections protected by the cache module's main lock into atomic blocks, and decorating required functions as `tm_callable`. The module had used atomic instructions for some memory accesses, and these were converted to full transactions in atomic blocks, so that collisions with these accesses will be detected.

4.3 Defining atomic blocks

Some transactions, after conversion contained code that belonged with the transactions, but didn't need necessarily to run atomically with the transaction. An example might be a transaction removing an object from the cache, and freeing its memory. While the removal operation must be protected inside a transaction, as it is using the shared memory structure of the cache, the memory release can happen any time later, since no other thread can point to the removed object after it had been removed from the cache.

For lock based systems, having the memory release as part of the critical section might cause a thread to hold the critical section a little longer than needed, but doesn't cause any problems other than that. On transactional memory systems, having accesses to other memory structures such as those required by memory management might cause collisions with other threads, thus slowing down the system in a similar way. In addition, the cleanup functions need to be transactified, which requires additional work both from the programmer and the compiler.

In our case, we chose not to transactify such functions, but instead remove them from the atomic section, and execute them after the transaction had committed. Although this requires some changes to the code, the changes are limited to the call-site, and need not modify any of the called libraries.

For example, the following critical section in the `open_entity` function (Figure 2) is responsible for retrieving a page to fulfill a request from the server. It will increment the reference count on the cached page, and register a decrement function to be invoked upon completing the request. When we converted the critical section into a transaction, we didn't want the function `apr_pool_cleanup_register` to be called from inside the transaction, as transactifying it would require working another library, the Apache Portable Runtime library, thus breaking encapsulation.

The semantics of requests and subrequest in apache guaranteed the request couldn't be completed before the return of this function, therefore we could move the registration of the cleanup function out of the atomic section, as seen in Figure 3. However, having a commit handler construct in the language would make such conversions easier, with it we could register a commit handler from within the transaction, and have the STM

system automatically invoke it once the transaction had committed.

4.4 Commit Handlers

Commit and undo handlers are pieces of code that are scheduled by a transaction to run when the transaction will commit, or abort, respectively. This mechanism, was suggested in [14]. Commit handlers are described there as a mechanism that allows finalization of tasks, for instance, a transactional system call such as write to file might have its permanent side effects be executed in a commit handler. Abort handlers are called when a transaction is aborted and can reverse the side-effects of a transaction. These handlers can sometimes be used to implement more efficient transactions. For example, if allocating memory inside a transaction, (assuming without a specialized memory allocator which is available in many STMs), the STM would need to log all the memory accesses to the memory management data structures, and undo these writes in case of an abort. A more efficient solution could be allocating the memory immediately, and in case of an abort just free the memory in an abort handler.

From our perspective, commit handlers could have been used to make the modifications we wanted in the atomic blocks, and move finalization functions out of atomic blocks just by registering them as commit handlers. In the given example, the call to `apr_pool_cleanup_register` could have been converted into a call registering this function as a commit handler.

5. Wish List

5.1 Handler Closures

While commit handlers can aid a lot in the process of transactifying a legacy application, their current syntax in Intel STM Compiler is very limiting. Handlers must be given as a pointer to function of a specific signature, so a developer trying to move a piece of code out of an atomic block, would still need to write a new function. It would be nice to have a language construct that defines a new commit handler right where it is being registered, however this requires the language to support closures, and of course there are many problems implementing those in a statically-typed language without garbage collection such as C.

```

static int open_entity(cache_handle_t *h, request_rec *r, const char *key) {
    ...
    if (sconf->lock) apr_thread_mutex_lock(sconf->lock);

    obj = (cache_object_t *) cache_find(sconf->cache_cache, key);
    if (obj) {
        if (obj->complete) {
            request_rec *rmain=r, *rtmp;
            apr_atomic_inc32(&obj->refcount);
            /* cache is worried about overall counts, not 'open' ones */
            cache_update(sconf->cache_cache, obj);

            /* If this is a subrequest, register the cleanup against the main
             * request. This will prevent the cache object from being cleaned up
             * from under the request after the subrequest is destroyed. */
            rtmp = r;
            while (rtmp) {
                rmain = rtmp;
                rtmp = rmain->main;
            }
            apr_pool_cleanup_register(rmain->pool, obj, decrement_refcount,
                apr_pool_cleanup_null);
        }
        else obj = NULL;
    }

    if (sconf->lock) apr_thread_mutex_unlock(sconf->lock);
    ...
}

```

Figure 2. Original `open_entity` function.

5.2 Statistics and Profiling

Intel's STM manager collects statistics about the transactions being run, their size, abort rates, etc. Unfortunately however, it cannot work with a multiprocess application such as Apache. This limits the ability to investigate the performance of converted applications to only limited runs with only one process, or having only black box measurements of the system.

6. Evaluation

6.1 Methodology

The transactified web server was evaluated using *Siege* [11], an HTTP load testing tool. The server was loaded with the set of UNIX man-pages - a set of small textual files typical of some web sites. Each page was served

using the *man2html* program, uncompressed and converted into HTML, to make sure the serving of files requires enough computational resources to make the use of cache worthwhile.

The *man2html* program is a Common Gateway Interface (CGI) program that serves unix manual (man) pages on internet sites. The pages are usually stores compressed in gzip format, and formatted using the troff format. The program receives a request for a man page from the webserver, uncompresses the required file and converts it to HTML. As every CGI program it outputs the result with relevant HTTP headers.

The default caching policy of apache forbids caching dynamically generated pages such as those of *man2html*, unless the HTTP headers of the resulting page clearly specify otherwise. To make caching of the *man2html*

```

static int open_entity(cache_handle_t *h, request_rec *r, const char *key)
{
    ...
    __tm_atomic {
        obj = (cache_object_t *) cache_find(sconf->cache_cache, key);
        if (obj) {
            if (obj->complete) {
                ++obj->refcount;
                /* cache is worried about overall counts, not 'open' ones */
                cache_update(sconf->cache_cache, obj);
            }
            else obj = NULL;
        }
    }

    /* Register the object for removal from the cache after cleanup */
    if (obj && obj->complete) {
        request_rec *rmain=r, *rtmp;
        /* If this is a subrequest, register the cleanup against the main
         * request. This will prevent the cache object from being cleaned up
         * from under the request after the subrequest is destroyed. */
        rtmp = r;
        while (rtmp) {
            rmain = rtmp;
            rtmp = rmain->main;
        }
        apr_pool_cleanup_register(rmain->pool, obj, decrement_refcount,
            apr_pool_cleanup_null);
    }
    ...
}

```

Figure 3. Transactified open_entity function.

pages possible, we modified man2html to output such headers, specifying the output can be cached for one hour.

The pages were requested randomly according to Zipf distribution, whose parameter s determines how frequently the most popular pages were visited, thus controlling the amount of locality in the requests.

The experiments were done with two computers connected with Gigabit ethernet. Both machines were an 8 processors SMP of quad core 2.3GHz AMD Opterons with 126GB of RAM, with a total of 32 cores.

6.2 Results

We compared the average latency and request throughput when running on different number of cores, and with different values. For every graph there are three experiments comparing the results of an Apache server running without a cache, a cached version without our transactional modifications, and the transactified version.

With low locality ($s = 0.1$), the cache has almost no benefit, and the penalty of the STM increases with the number of processors, causing it to be less practical, as we learn from Figure 5 and 6.

```

static int open_entity(cache_handle_t *h, request_rec *r, const char *key) {
    ...
    __tm_atomic {
        obj = (cache_object_t *) cache_find(sconf->cache_cache, key);
        if (obj) {
            if (obj->complete) {
                request_rec *rmain=r, *rtmp;
                ++obj->refcount;
                /* cache is worried about overall counts, not 'open' ones */
                cache_update(sconf->cache_cache, obj);

                /* If this is a subrequest, register the cleanup against the main
                 * request. This will prevent the cache object from being cleaned
                 * up from under the request after the subrequest is destroyed. */
                rtmp = r;
                while (rtmp) {
                    rmain = rtmp;
                    rtmp = rmain->main;
                }
                on_commit {
                    apr_pool_cleanup_register(rmain->pool, obj, decrement_refcount,
                        apr_pool_cleanup_null);
                }
            }
            else obj = NULL;
        }
    }
    ...
}

```

Figure 4. open_entity function with a commit handler closure

With higher locality, such as $s = 1$ (Figure 7, 8), the overhead of STM is still significant, but using the cache is now worthwhile, even with STM. It is worth noting that when running on a single machine, without the overhead of the network (Figure 9, 10), the decrease in performance caused by STM was even less significant. [XX Why? XX]

With even higher locality ($s = 2$), the contention on the cache is high even with small number of cores, and both versions show similar results. (Figure 11, 12).

As notes in Section 5.2, Intel STM manager statistics mechanism was unable to work with multiple processes. In order to collect internal data about the transactions we ran Apache in a single process debugging mode. In Figure 13, the results show that STM achieves comparable results to the original version. In Figure 14

the average number of aborts per second is shown, for the experiment with the transactified version. The increase in contention shown here explains the decrease in request throughput.

7. Conclusions

- Out of 340,000 lines of code in the Apache Web server, the cache module is comprised of only 6651 lines of code, of which only 273 lines were changed by our modifications. This shows the importance of being able to modify only encapsulated sections of the code, and interoperating with legacy code that still uses locks.
- Having commit handlers in the STM system is not only needed for creating efficient open transactions,




Figure 5. Request throughput, $s = 0.1$: Low locality. The penalty of STM increases to the point it outweighs the benefit of using the cache.

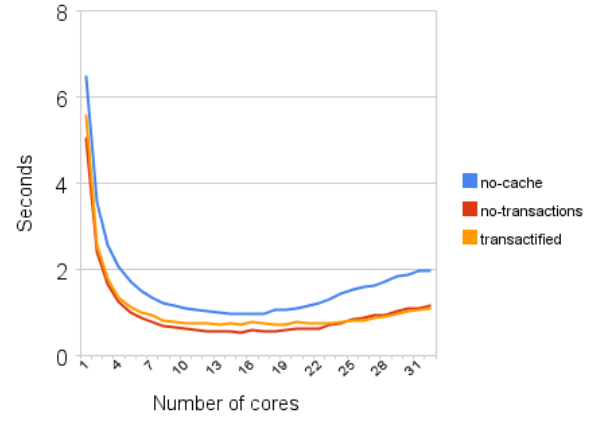


Figure 8. Average response time, $s = 1$: Medium locality




Figure 6. Average response time, $s = 0.1$: Low locality




Figure 9. Request throughput, Single host, $s = 1$: Medium locality. With a single host, the additional latency of the network is eliminated, achieving generally higher results than with two hosts.




Figure 7. Request throughput, $s = 1$: Medium locality. The penalty of STM remains but the cache is an improvement.




Figure 10. Average response time, Single host, $s = 1$: Medium locality




Figure 11. Request throughput, $s = 2$: High locality.




Figure 12. Average response time, $s = 2$: High locality




Figure 13. Request throughput, $s = 1$: Medium locality, a single process. Without the additional processes apache fails to scale with high number of cores. STM suffers a decrease in performance but is comparable.




Figure 14. Abort Rate, $s = 1$: Medium locality, a single process. The abort rate increase with the number of cores might explain the decrease in transaction rate.

but can also aid the process of transactifying legacy code.

- There is great importance in working on real-world applications as they may reveal challenges resulting from engineering problems and not only algorithmic and theoretical problems.

There are many STM systems currently available, and one immediate direction would be to compare them using this benchmark. This would require writing plugins for any such system to match Intel's TM ABI.

In addition, there are other applications that might be interesting as transactional memory applications, following the methods we used.

Acknowledgments

We would like to thank Torvald Riegel from the Transactional Memory Research group at Dresden University of Technology, for helping us with our efforts to use TANGER for our research.

References

- [1] Albert Cohen. GCC for Transactional Memory. HiPEAC, 7 June 2008. URL <http://www.hipeac.net/node/2419>.
- [2] Woongki Baek, Chi Cao Minh, Martin Trautmann, Christos Kozyrakis, and Kunle Olukotun. The opentm transactional application programming interface. In *Proceedings of the 16th International Conference on Parallel Architectures and Compilation Techniques*, pages 376–387. Sep 2007.

- [3] Chi Cao Minh, JaeWoong Chung, Christos Kozyrakis, and Kunle Olukotun. Stamp: Stanford transactional applications for multi-processing. In *IISWC '08: Proceedings of The IEEE International Symposium on Workload Characterization*, September 2008.
- [4] O. Shalev D. Dice and N. Shavit. Transactional locking ii. In *Proc. of the 20th International Symposium on Distributed Computing (DISC 2006)*, pages 194–208, 2006.
- [5] Pascal Felber, Christof Fetzer, Ulrich Müller, Torvald Riegel, Martin Süßkraut, and Heiko Sturzrehm. Transactifying applications using an open compiler framework. In *TRANSACT*, August 2007.
- [6] Pascal Felber, Christof Fetzer, and Torvald Riegel. Dynamic performance tuning of word-based software transactional memory. In *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, 2008.
- [7] Keir Fraser. *Practical lock freedom*. PhD thesis, Cambridge University Computer Laboratory, 2003. Also available as Technical Report UCAM-CL-TR-579.
- [8] Rachid Guerraoui, Michal Kapalka, and Jan Vitek. STMBench7: A Benchmark for Software Transactional Memory. In *Second European Systems Conference EuroSys2007*, 2007.
- [9] Maurice Herlihy, Victor Luchangco, Mark Moir, and William N. Scherer, III. Software transactional memory for dynamic-sized data structures. In *PODC '03: Proceedings of the twenty-second annual symposium on Principles of distributed computing*, pages 92–101, New York, NY, USA, 2003. ACM. ISBN 1-58113-708-7. doi: <http://doi.acm.org/10.1145/872035.872048>.
- [10] Intel. Transactional memory ABI. URL <http://software.intel.com/file/8097>.
- [11] Jeffrey Fulmer. Siege HTTP regression testing and benchmarking utility. URL <http://www.joedog.org/JoeDog/Siege>.
- [12] Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the 2004 International Symposium on Code Generation and Optimization (CGO'04)*, Palo Alto, California, Mar 2004.
- [13] Virendra J. Marathe, Michael F. Spear, Christopher Heriot, Athul Acharya, David Eisenstat, William N. Scherer Iii, and Michael L. Scott. Lowering the overhead of nonblocking software transactional memory. In *Dept. of Computer Science, Univ. of Rochester*, 2006.
- [14] Austen Mcdonald, Jaewoong Chung, Brian D. Carlstrom, Chi Cao Minh, Hassan Chafi, Christos Kozyrakis, and Kunle Olukotun. Architectural semantics for practical transactional memory. In *In Proceedings of the 33rd International Symposium on Computer Architecture*, pages 53–65. IEEE Computer Society, 2006.
- [15] Yang Ni, Adam Welc, Ali-Reza Adl-Tabatabai, Moshe Bach, Sion Berkowits, James Cownie, Robert Geva, Sergey Kozhukow, Ravi Narayanaswamy, Jeffrey Olivier, Serguei Preis, Bratin Saha, Ady Tal, and Xinmin Tian. Design and implementation of transactional constructs for c/c++. In *OOPSLA '08: Proceedings of the 23rd ACM SIGPLAN conference on Object oriented programming systems languages and applications*, pages 195–212, New York, NY, USA, 2008. ACM. ISBN 978-1-60558-215-3. doi: <http://doi.acm.org/10.1145/1449764.1449780>.
- [16] The Apache Software Foundation. Apache HTTP Server Project. URL <http://httpd.apache.org/>.