Transactifying Apache

Haggai Eran   Ohad Lutzky   Zvika Guz   Idit Keidar

Technion

haggaie@tx.technion.ac.il   lutzky@gmail.com   zguz@tx.technion.ac.il   idish@ee.technion.ac.il

Abstract

Apache is a large-scale industrial multi-process and multi-threaded application, which uses lock-based synchronization. We report on our experience in modifying Apache to employ transactional memory instead of locks, a process we refer to as *transactification*; we are not aware of any previous efforts to transactify legacy software of such a large scale. Along the way, we learned some valuable lessons about which tools one should use, which parts of the code one should transactify and which are better left untouched, as well as on the intricacy of commit blocks. We also stumbled across weaknesses of existing software transactional memory (STM) toolkits, leading us to identify desirable features they are currently lacking. Finally, we present performance results from running Apache on a 32-core machine, showing that, surprisingly, the performance of the STM-based version is close to that of the lock-based version. These results suggest that there are applications for which the overhead of using a software-only implementation of transactional memory is insignificant.

Categories and Subject Descriptors   CR-number [*subcategory*]: third-level

General Terms   term1, term2

Keywords   keyword1, keyword2

1. Introduction

The text of the paper begins here.

Figure 1. Apache worker MPM architecture.

2.   Related Work

3.   Background: Software and Tools

3.1   Apache

Apache (The Apache Software Foundation a) HTTP server is a popular web server application written in C. It supports working on multiprocessor machines with several multi-processing modules (MPMs) each offering a different strategy for handling requests and distributing the work. The most popular threaded MPM is the *worker* MPM, which works by running multiple worker-threads under several processes, each thread handles a single request at a time. In each such process there are several worker threads, and also a listener thread that fetches incoming requests and dispatches them to the available workers, as illustrated in Figure 1.

There are not many points of interaction between the worker threads themselves, where transactional memory can be used. One such place is Apache's memory cache implemented by the `mod_mem_cache` (The Apache Software Foundation d) module. This module

enables the workers of each process to share a cache of recently served requests. A new request can be served from the memory cache, and save the time required to access the disk and generate the requested page. Since the cache is shared between multiple threads, it is synchronized by a single lock, therefore a good candidate for converting into transactional memory.

Apache's cache is implemented with a couple of modules. The first, `mod_cache` (The Apache Software Foundation b), implements the logic related to caching. It tests the metadata of each requests to see if it can be supplied from the cache, according to the request's HTTP headers and the system configuration. It uses one of the underlying cache implementation modules, `mod_mem_cache` or `mod_disk_cache` (The Apache Software Foundation c) to do the actual caching.

The `mod_mem_cache` module implements a memory cache using a shared hash table and priority queue. The key to the hash table is the URL of the request, converted into a canonical form. The cache is limited both by size and by the number of elements, and by memory size, so on insertion, sometimes lower priority entries are removed from the cache. The priority is deterimined by one of two algorithms: LRU, removing the least recently used entries first, and GDSF (Greedy Dual Size Frequency) assigning score to entries based on the cost of a cache miss, and the entry size.

3.2   C STM Systems

C and C++ STM systems divide into two kinds: Library based and compiler based. Library based STMs are built as a C library. Every transaction begins with a call into the library, and commits by another call. All reads and writes to global variables must be done through special library functions when in a transaction. This requires a great amount of work for converting an application to use STM. Not only accesses to global memory in the function that started the transaction must be converted, but also any access from any function being called from this function.

In contrast, compiler-based STM use a specialized compiler, which has extended syntax for transactional memory atomic blocks. The compiler can then automatically convert memory accesses inside transactions into calls to the underlying library, a process sometimes referred as *transactification*.

3.2.1   TANGER

The TANGER (Felber et al. 2007) transctifying compiler is an open-source academic compiler extension for LLVM (Lattner and Adve 2004), an extensible compiler framework. Tanger aims at creating a transactifying compiler that is independant of the STM system used. It works with the tinySTM (Felber et al. 2008) library, but can easily be extended to use other STM libraries by writing a simple plugin.

TANGER is accompanied with the TARIFA (Felber et al. 2007) tool, which transactifies compiled binaries, even without the sources, which might be a very important advantage when modifying legacy code.

3.2.2   Intel STM Compiler

Intel has published (Ni et al. 2008) an experimental STM compiler based on their industrial compiler ICC. It solved the above problem by adding some new function attributes to the language that tell the compiler which functions need to be transactified. The attribute `tm_callable` tells the compiler that a transactified version of the function will be needed. This way only functions that are required inside transactions can be marked as `tm_callable` and be transactified.

Although ICC uses a proprietery STM manager, Intel has published their ABI (Intel) allowing for other STM managers to replace their own. This feature and its selective transactification ability were the main reasons why we preferred ICC.

In the latest version of Intel STM Compiler, support for abort and commit handlers was in fact added to the system, by registering a callback function from inside a transaction.

An extension to the Gnu Compiler Collection (GCC) is being developed (Albert Cohen 2008) to enable transactional memory support for GCC. It is intended to work with tinySTM, but being open source, other STM systems will probably be ported too. The syntax of the C/C++ language extensions is designed to be compatible with ICC. This means that applications converted for ICC will probably be compilable under GCC with this extension, without much modification.

4.   Transactifycation

4.1   Which STM to use?

TANGER created a transactified version of each function in a compilation unit. Every function call inside a transaction was then converted to a call to the new ver-

sion. This method is a major disadvantage when working on a large application. Many functions do not need a transactified version and this causes uneeded work for the compiler and the linker. Moreover, sometimes the transactifyication might fail because of calls to functions whose source is not available and cannot be transactified. This can cause the entire build process to fail, where in fact the code can be transactified without any error.

This was the main reason why we eventually chose to use Intel's STM Compiler. Lately a new release of tanger was announced, one in which the developer can annotate which functions should be transactified, however we didn't get a chance to try it.

4.2 What to transactify

The conversion process included converting critical sections protected by the cache module's main lock into atomic blocks, and decorating required functions as `tm_callable`. The module had used atomic instructions for some memory accesses, and these were converted to full transactions in atomic blocks, so that collisions with these accesses will be detected.

4.3 Defining atomic blocks

Some transactions, after conversion contained code that belonged with the transactions, but didn't need neccessarily to run atomically with the transaction. An example might be a transaction removing an object from the cache, and freeing its memory. While the removal operation must be protected inside a transaction, as it is using the shared memory structure of the cache, the memory release can happen any time later, since no other thread can point to the removed object after it had been removed from the cache.

For lock based systems, having the memory release as part of the critical section might cause a thread to hold the critical section a little longer than needed, but doesn't cause any problems other than that. On transactional memory systems, having accesses to other memory structures such as those required by memory management might cause collisions with other threads, thus slowing down the system in a similar way. In addition, the cleanup functions need to be transactified, which requires additional work both from the programmer and the compiler.

In our case, we chose not to transactify such functions, but instead remove them from the atomic section, and execute them after the transaction had committed.

Although this requires some changes to the code, the changes are limited to the call-site, and need not modify any of the called libraries.

For example, the following critical section in the `open_entity` function (Figure 2) is responsible for retrieving a page to fullfill a request from the server. It will increment the reference count on the cached page, and register a decrement function to be invoked upon completing the request. When we converted the critical section into a transaction, we didn't want the function `apr_pool_cleanup_register` to be called from inside the transaction, as transactifying it would require working an another library, the Apache Portable Runtime library, thus breaking encapsulation.

The semantics of requests and subrequest in apache guranteed the request couldn't be completed before the return of this function, therefore we could move the registration of the cleanup function out of the atomic section, as seen in Figure 3. However, having a commit handler construct in the language would make such conversions easier, with it we could register a commit handler from within the transaction, and have the STM system automatically invoke it once the transaction had committed.

4.4 Commit Handlers

Commit and undo handlers are pieces of code that are scheduled by a transaction to run when the transaction will commit, or abort, respectively. This mechanism, was suggested in (Mcdonald et al. 2006). Commit handlers are described there as a mechanism that allows finalization of tasks, for instance, a transactional system call such as write to file might have its permanent side effects be executed in a commit handler. Abort handlers are called when a transaction is aborted and can reverse the side-effects of a transaction. These handlers can sometimes be used to implement more efficient transactions. For example, if allocating memory inside a transaction, (assuming without a specialized memory allocator which is available in many STMs), the STM would need to log all the memory accesses to the memory management data structures, and undo these writes in case of an abort. A more efficient solution could be allocating the memory immediately, and in case of an abort just free the memory in an abort handler.

From our perspective, commit handlers could have been used to make the modifications we wanted in the

```
static int open_entity(cache_handle_t *h, request_rec *r, const char *key) {
  ...
  if (sconf->lock) apr_thread_mutex_lock(sconf->lock);

  obj = (cache_object_t *) cache_find(sconf->cache_cache, key);
  if (obj) {
    if (obj->complete) {
      request_rec *rmain=r, *rtmp;
      apr_atomic_inc32(&obj->refcount);
      /* cache is worried about overall counts, not 'open' ones */
      cache_update(sconf->cache_cache, obj);

      /* If this is a subrequest, register the cleanup against the main request.
       * This will prevent the cache object from being cleaned up from under the
       * request after the subrequest is destroyed. */
      rtmp = r;
      while (rtmp) {
        rmain = rtmp;
        rtmp = rmain->main;
      }
      apr_pool_cleanup_register(rmain->pool, obj, decrement_refcount,
        apr_pool_cleanup_null);
    }
    else obj = NULL;
  }

  if (sconf->lock) apr_thread_mutex_unlock(sconf->lock);
  ...
}
```

Figure 2. Original `open_entity` function.

atomic blocks, and move finalization functions out of atomic blocks just by registering them as commit handlers. In the given example, the call to `apr_pool_cleanup_register` could have been converted into a call registering this function as a commit handler.

fines a new commit handler right where it is being registered, however this requires the language to support closures, and of course there are many problems implementing those in a statically-typed language without garbage collection such as C.

5. Wish List

5.1 Handler Closures

While commit handlers can aid a lot in the process of transactifying a legacy application, their current syntax in Intel STM Compiler is very limiting. Handlers must be given as a pointer to function of a specific signature, so a developer trying to move a piece of code out of an atomic block, would still need to write a new function. It would be nice to have a language construct that de-

5.2 Statistics and Profiling

Intel's STM manager collects statistics about the transactions being run, their size, abort rates, etc. Unfortunately however, it cannot work with a multiprocess application such as Apache. This limits the ability to investigate the performance of converted applications to only limited runs with only one process, or having only black box measurements of the system.

```c
static int open_entity(cache_handle_t *h, request_rec *r, const char *key)
{
  ...
  __tm_atomic {
    obj = (cache_object_t *) cache_find(sconf->cache_cache, key);
    if (obj) {
      if (obj->complete) {
        ++obj->refcount;
        /* cache is worried about overall counts, not 'open' ones */
        cache_update(sconf->cache_cache, obj);
      }
      else obj = NULL;
    }
  }

  /* Register the object for removal from the cache after cleanup */
  if (obj && obj->complete) {
    request_rec *rmain=r, *rtmp;
    /* If this is a subrequest, register the cleanup against the main request.
     * This will prevent the cache object from being cleaned up from under the
     * request after the subrequest is destroyed.
     */
    rtmp = r;
    while (rtmp) {
      rmain = rtmp;
      rtmp = rmain->main;
    }
    apr_pool_cleanup_register(rmain->pool, obj, decrement_refcount,
      apr_pool_cleanup_null);
  }
  ...
}
```

Figure 3. Transactified `open_entity` function.

6. Evaluation

6.1 Methodology

The transactified web server was evaluated using *Siege* (Jeffrey Fulmer), an HTTP load testing tool. The server was loaded with the set of UNIX man-pages - a set of small textual files typical of some web sites. Each page was served using the *man2html* (man) program, uncompressed and converted into HTML, to make sure the serving of files requires enough computational resources to make the use of cache worthwhile.

The man2html program is a Common Gateway Interface (CGI) program that serves unix manual (man) pages on internet sites. The pages are usually stores compressed in gzip format, and formatted using the troff format. The program receives a request for a man page from the webserver, uncompresses the required file and converts it to HTML. As every CGI program it outputs the result with relevant HTTP headers.

The default caching policy of apache forbids caching dynamically generated pages such as those of man2html, unless the HTTP headers of the resulting page clearly specify otherwise. To make caching of the man2html pages possible, we modified man2html to output such headers, specifying the output can be cached for one hour.

```
static int open_entity ( cache_handle_t *h, request_rec *r, const char *key) {
    ...
    __tm_atomic {
        obj = (cache_object_t *) cache_find ( sconf->cache_cache, key);
        if (obj) {
            if (obj->complete) {
                request_rec *rmain=r, *rtmp;
                ++obj->refcount;
                /* cache is worried about overall counts, not 'open' ones */
                cache_update ( sconf->cache_cache, obj);

                /* If this is a subrequest, register the cleanup against the main request.
                 * This will prevent the cache object from being cleaned up from under the
                 * request after the subrequest is destroyed. */
                rtmp = r;
                while (rtmp) {
                    rmain = rtmp;
                    rtmp = rmain->main;
                }
                on_commit {
                    apr_pool_cleanup_register ( rmain->pool, obj, decrement_refcount,
                        apr_pool_cleanup_null );
                }
            }
        }
        else obj = NULL;
        }
    }
    ...
}
```

Figure 4. `open_entity` function with a commit handler closure

The pages were requested randomly according to Zipf distribution, whose parameter s determines how frequently the most popular pages were visited, thus controlling the amount of locality in the requests.

The experiments were done with two computers connected using Gigabit ethernet. The machine running the server was a 4 processors SMP of dual core 2.66GHz Xeons with 8GB of RAM, and the client machine was a 2 processors SMP of quad core 2.33Ghz E5410 Xeons with 8GB of RAM. [TBD: Update to neo and trinity]

6.2 Results

We compared the average latency and request throughput when running on different number of cores, and with different values. For every graph there are three experiments comparing the results of an Apache server running without a cache, a cached version without our transactional modifications, and the transactified version.

7. Conclusions

- Out of 340,000 lines of code in the Apache Web server, the cache module is comprised of only 6651 lines of code, of which only 273 lines were changed by our modifications. This shows the importance of being able to modify only encapsulated sections of the code, and interoperating with legacy code that still uses locks.

- Having commit handlers in the STM system is not only needed for creating efficient open transactions,

**Figure 5.** Transaction Rate, Single host, $s = 1$

Figure 6. Average Response Time, Single host, $s = 1$

Figure 7. Average Response Time, $s = 1$

Figure 8. Average Response Time, $s = 2$

Figure 9. Transaction Rate, $s = 1$

**Figure 10.** Transaction Rate, $s = 2$

**Figure 11.** Transaction Rate, $s = 1$, a single process

Figure 12. Abort Rate, $s = 1$, a single process

but can also aid the process of transactifying legacy code.

- There is great importance in working on real-world applications as they may reveal challenges resulting from engineering problems and not only algorithmic and theoretical problems.

There are many STM systems currently available, and one immediate direction would be to compare them using this benchmark. This would require writing plugins for any such system to match Intel's TM ABI.

In addition, there are other applications that might be interesting as transactional memory applications, following the methods we used.

A. Appendix Title

This is the text of the appendix, if you need one.

Acknowledgments

Acknowledgments, if needed.

References

Man2html package source code. URL http://packages.ubuntu.com/source/intrepid/man2html.

Albert Cohen. GCC for Transactional Memory. HiPEAC, 7 June 2008. URL http://www.hipeac.net/node/2419.

Pascal Felber, Christof Fetzer, Ulrich Müller, Torvald Riegel, Martin Süßkraut, and Heiko Sturzrehm. Transactifying applications using an open compiler framework. In *TRANSACT*, August 2007.

Pascal Felber, Christof Fetzer, and Torvald Riegel. Dynamic performance tuning of word-based software transactional memory. In *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, 2008.

Intel. Transactional memory ABI. URL http://software.intel.com/file/8097.

Jeffrey Fulmer. Siege HTTP regression testing and benchmarking utility. URL http://www.joedog.org/JoeDog/Siege.

Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the 2004 International Symposium on Code Generation and Optimization (CGO'04)*, Palo Alto, California, Mar 2004.

A. Mcdonald, J. Chung, B. Carlstrom, C. Minh, H. Chafi, C. Kozyrakis, and K. Olukotun. Architectural semantics for practical transactional memory, 2006. URL http://citeseer.ist.psu.edu/mcdonald06architectural.html.

Yang Ni, Adam Welc, Ali-Reza Adl-Tabatabai, Moshe Bach, Sion Berkowits, James Cownie, Robert Geva, Sergey Kozhukow, Ravi Narayanaswamy, Jeffrey Olivier, Serguei Preis, Bratin Saha, Ady Tal, and Xinmin Tian. Design and implementation of transactional constructs for c/c++. In *OOPSLA '08: Proceedings of the 23rd ACM SIGPLAN conference on Object oriented programming systems languages and applications*, pages 195–212, New York, NY, USA, 2008. ACM. ISBN 978-1-60558-215-3. doi: http://doi.acm.org/10.1145/1449764.1449780.

The Apache Software Foundation. Apache HTTP Server Project, a. URL http://httpd.apache.org/.

The Apache Software Foundation. Apache's mod_cache Memory Cache Module, b. URL http://httpd.apache.org/docs/2.2/mod/mod_cache.html.

The Apache Software Foundation. Apache's mod_disk_cache Memory Cache Module, c. URL http://httpd.apache.org/docs/2.2/mod/mod_disk_cache.html.

The Apache Software Foundation. Apache's mod_mem_cache Memory Cache Module, d. URL http://httpd.apache.org/docs/2.2/mod/mod_mem_cache.html